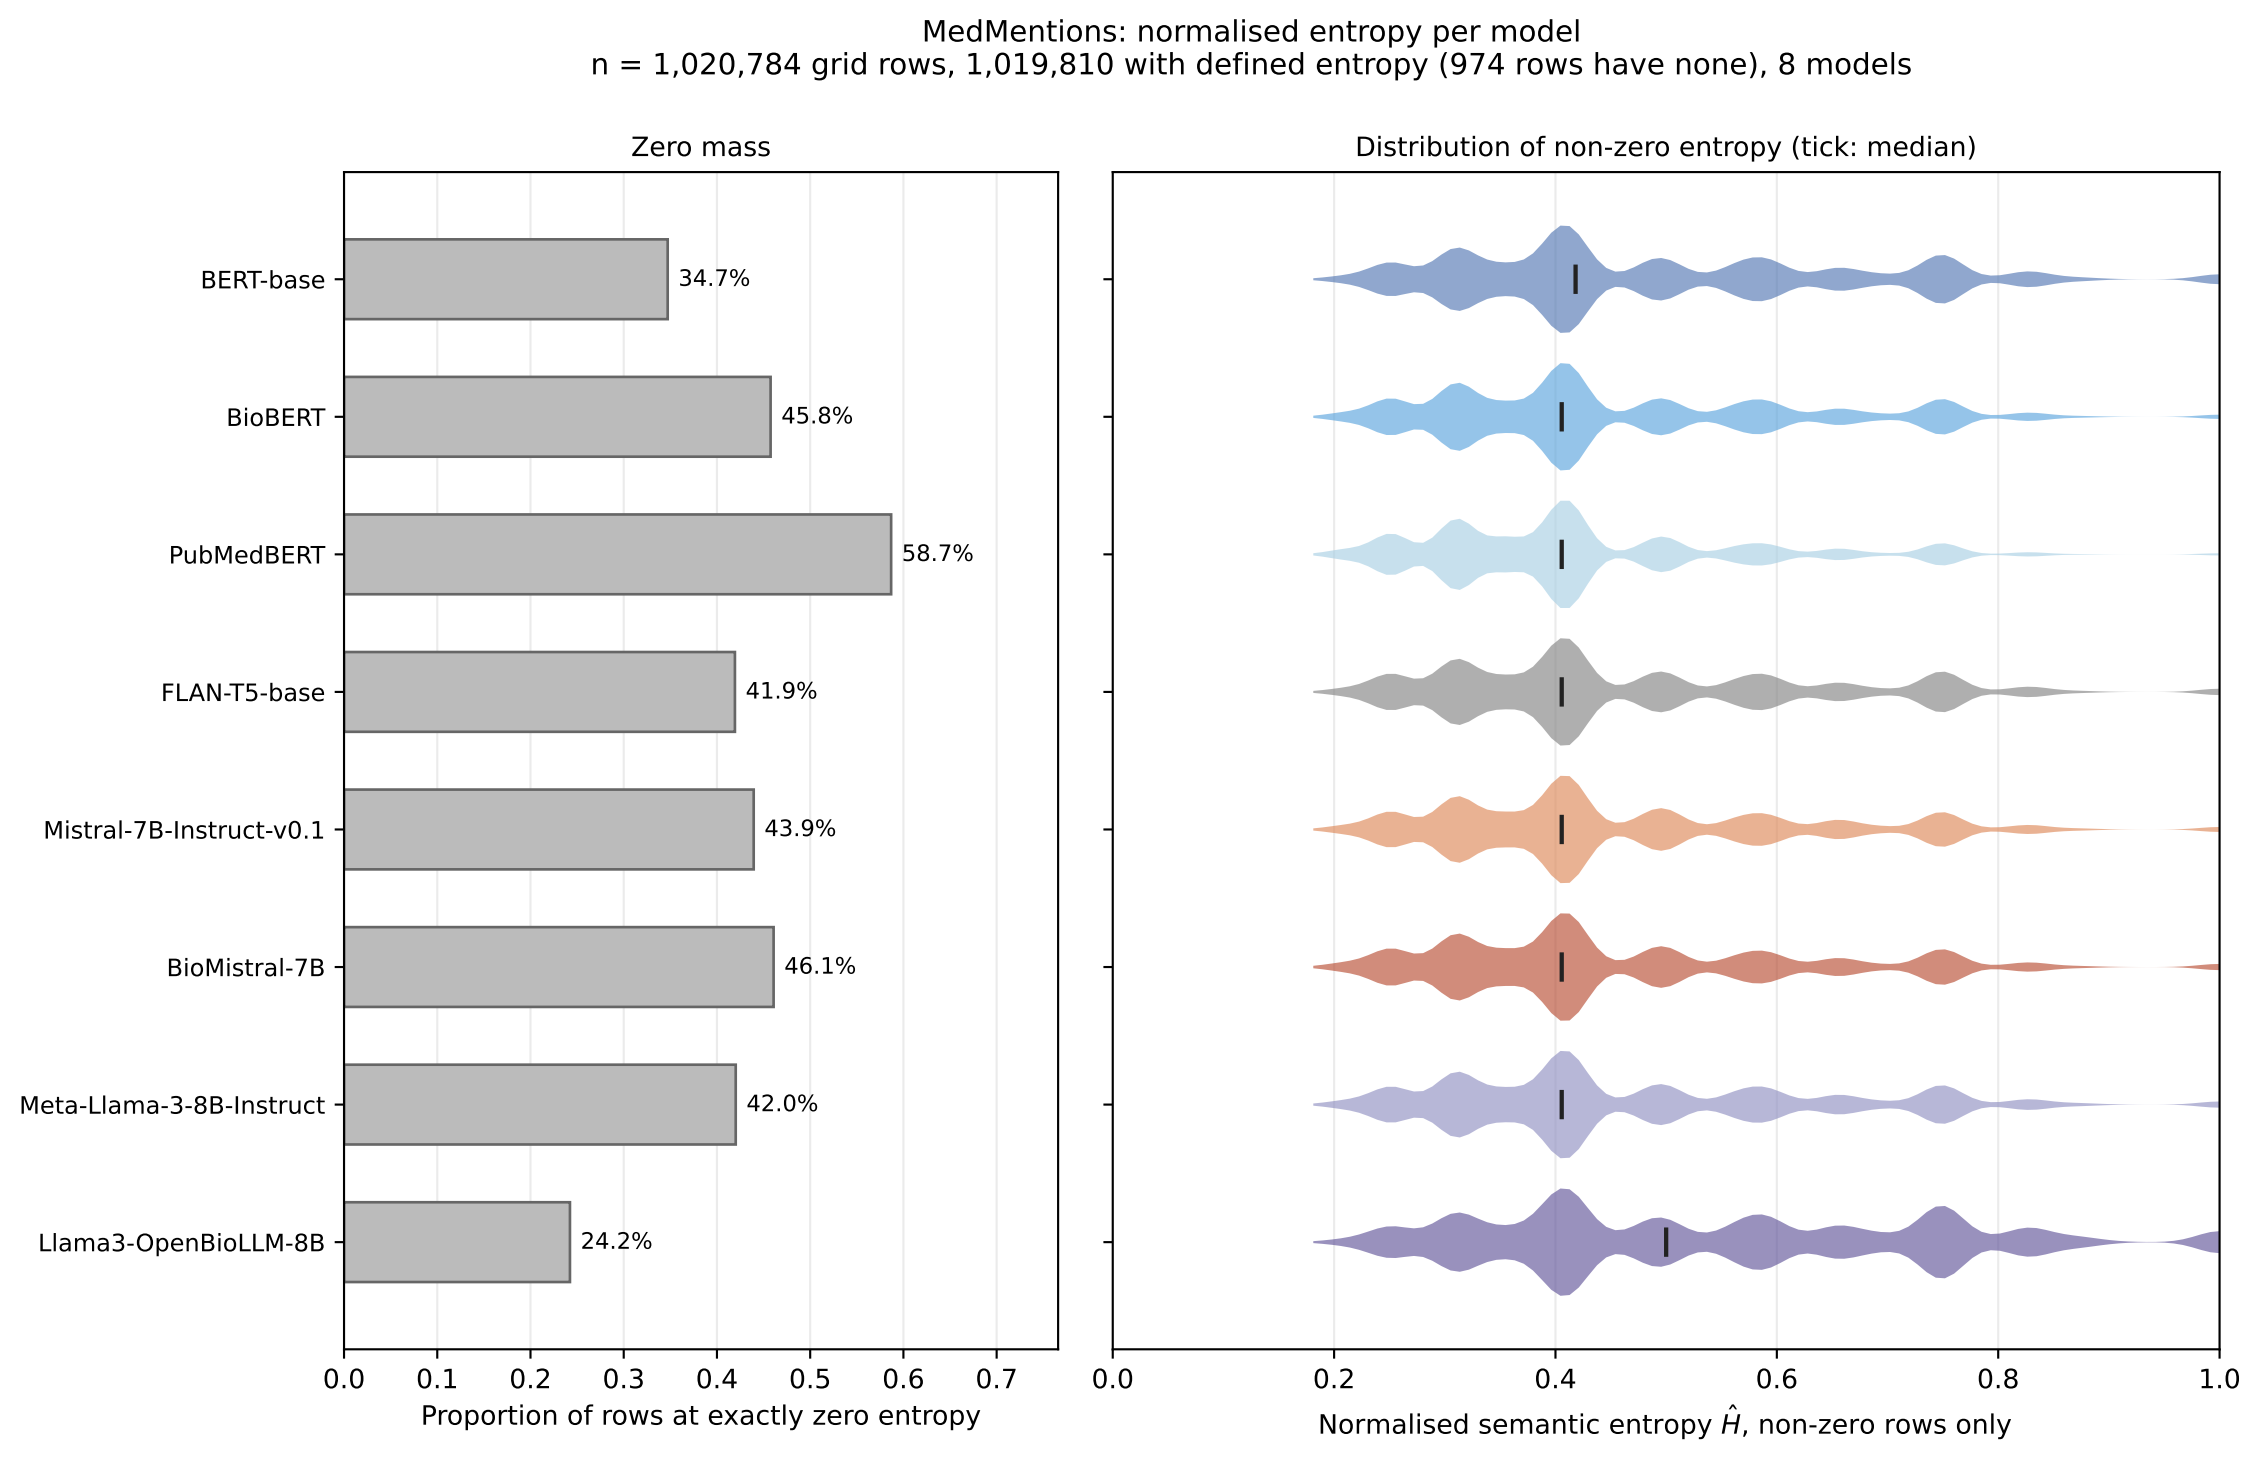
MedMentions: normalised entropy per model
n = 1,020,784 grid rows, 1,019,810 with defined entropy (974 rows have none), 8 models
Zero mass
BERT-base
34.7%
BioBERT
45.8%
PubMedBERT
58.7%
FLAN-T5-base
41.9%
Mistral-7B-Instruct-v0.1
43.9%
BioMistral-7B
46.1%
Meta-Llama-3-8B-Instruct
42.0%
Llama3-OpenBioLLM-8B
24.2%
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
Proportion of rows at exactly zero entropy
Distribution of non-zero entropy (tick: median)
0.0
0.2
0.4
0.6
0.8
1.0
Normalised semantic entropy $\hat{H}$, non-zero rows only