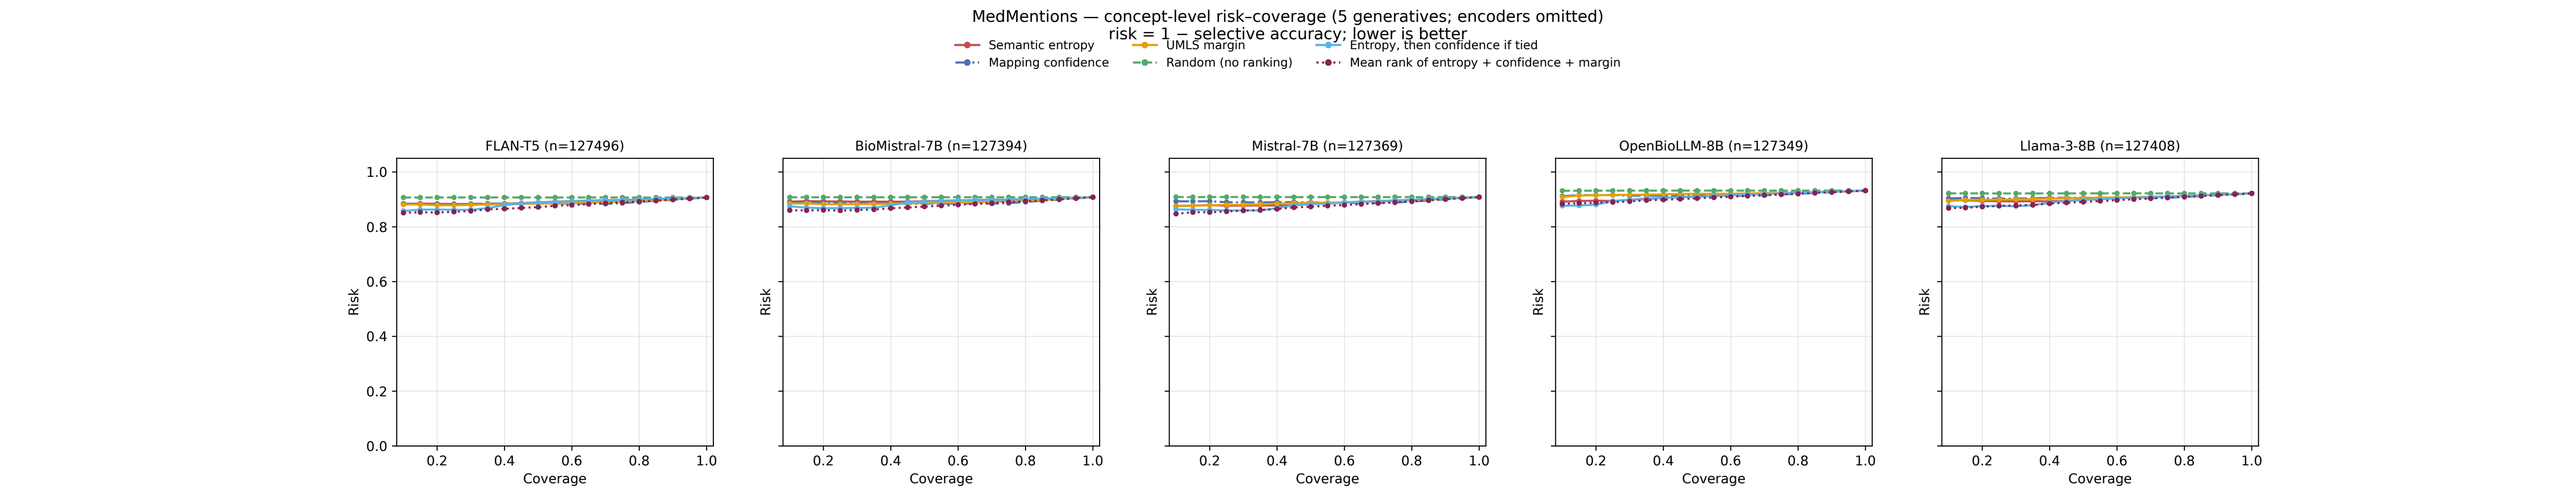

MedMentions encoders are OMITTED from this panel. They are direct-CUI: mapping confidence is constant 1.0 (a degenerate ranker) and UMLS margin is the input-mention-span gap (mention difficulty), not a model-uncertainty axis. Earlier versions of this figure carried forward pilot-vintage encoder curves and disclosed the mixed sample sizes in a footnote; the cut-off re-run quarantined those margins, so the panel is now 5 generatives at full scale and one vintage throughout.